

Practical No: 1

Aim: Document Indexing and Retrieval

- Implement an inverted index construction algorithm.
- Build a simple document retrieval system using the constructed index.

Code:

Input:

A) Implement an inverted index construction algorithm.

```
import nltk
# Import NLTK to download stopwords
from nltk.corpus import stopwords
# Import stopwords from NLTK

# Define the documents
document1 = "The quick brown fox jumped over the lazy dog"
document2 = "The lazy dog slept in the sun"

# Get the stopwords for English language from NLTK
nltk.download('stopwords')
stopWords = stopwords.words('english')

# Step 1: Tokenize the documents
# Convert each document to lowercase and split it into words
tokens1 = document1.lower().split()
tokens2 = document2.lower().split()

# Combine the tokens into a list of unique terms
terms = list(set(tokens1 + tokens2))

# Step 2: Build the inverted index
# Create an empty dictionary to store the inverted index as well as a dictionary
# to store number of occurrences
```

```
inverted_index = {}
occ_num_doc1 = {}
occ_num_doc2 = {}

# For each term, find the documents that contain it
for term in terms:
    if term in stopWords:
        continue
    documents = []
    if term in tokens1:
        documents.append("Document 1")
        occ_num_doc1[term] = tokens1.count(term)
    if term in tokens2:
        documents.append("Document 2")
        occ_num_doc2[term] = tokens2.count(term)

inverted_index[term] = documents
```

Step 3: Print the inverted index

```
for term, documents in inverted_index.items():
    print(term, "->", end=" ")
    for doc in documents:
        if doc == "Document 1":
            print(f"{doc} ({occ_num_doc1.get(term, 0)}),", end=" ")
        else:
            print(f"{doc} ({occ_num_doc2.get(term, 0)}),", end=" ")
    print()
print("Performed by TYCS")
```

Output:

```
[nltk_data] Downloading package stopwords to
[nltk_data]      C:\Users\deepa\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
quick -> Document 1 (1),
lazy -> Document 1 (1), Document 2 (1),
sun -> Document 2 (1),
jumped -> Document 1 (1),
fox -> Document 1 (1),
slept -> Document 2 (1),
dog -> Document 1 (1), Document 2 (1),
brown -> Document 1 (1),
Performed by 740_Pallavi & 743_Deepak
```

OR

```
# Define the documents
```

```
document1 = "The quick brown fox jumped over the lazy dog."
```

```
document2 = "The lazy dog slept in the sun."
```

Step 1: Tokenize the documents

```
# Convert each document to lowercase and split it into words
```

```
tokens1 = document1.lower().split()
```

```
tokens2 = document2.lower().split()
```

```
# Combine the tokens into a list of unique terms
```

```
terms = list(set(tokens1 + tokens2))
```

Step 2: Build the inverted index

```
# Create an empty dictionary to store the inverted index
```

```
inverted_index = { }
```

```
# For each term, find the documents that contain it
```

```
for term in terms:
```

```
    documents = []
```

```
    if term in tokens1:
```

```
        documents.append("Document 1")
```

```
    if term in tokens2:
```

```
        documents.append("Document 2")
```

```
    inverted_index[term] = documents
```

Step 3: Print the inverted index

```
for term, documents in inverted_index.items():  
    print(term, "->", ", ".join(documents))
```

Practical No: 2

Aim: Retrieval Models

- Implement the Boolean retrieval model and process queries.
- Implement the vector space model with TF-IDF weighting and cosine similarity.

Code:

Input:

A) Implement the Boolean retrieval model and process queries:

```
documents = {
    1: "apple banana orange",
    2: "apple banana",
    3: "banana orange",
    4: "apple"
}

# Function to build an inverted index using dictionaries
def build_index(docs):
    index = {} # Initialize an empty dictionary to store the inverted index
    for doc_id, text in docs.items(): # Iterate through each document and its text
        terms = set(text.split()) # Split the text into individual terms
        for term in terms: # Iterate through each term in the document
            if term not in index:
                index[term] = {doc_id} # If the term is not in the index, create a new
set with document ID
            else:
                index[term].add(doc_id) # If the term exists, add the document ID to
its set
    return index # Return the built inverted index
```

Building the inverted index

```
inverted_index = build_index(documents)
```

Function for Boolean AND operation using inverted index

```
def boolean_and(operands, index):
```

```
    if not operands: # If there are no operands, return all document IDs
```

```
        return list(range(1, len(documents) + 1))
```

```
    result = index.get(operands[0], set()) # Get the set of document IDs for the  
    first operand
```

```
    for term in operands[1:]: # Iterate through the rest of the operands
```

```
        result = result.intersection(index.get(term, set())) # Compute intersection  
    with sets of document IDs
```

```
    return list(result) # Return the resulting list of document IDs
```

Function for Boolean OR operation using inverted index

```
def boolean_or(operands, index):
```

```
    result = set() # Initialize an empty set to store the resulting document IDs
```

```
    for term in operands: # Iterate through each term in the query
```

```
        result = result.union(index.get(term, set())) # Union of sets of document  
    IDs for each term
```

```
    return list(result) # Return the resulting list of document IDs
```

Function for Boolean NOT operation using inverted index

```
def boolean_not(operand, index, total_docs):
```

```
    operand_set = set(index.get(operand, set())) # Get the set of document IDs  
    for the operand
```

```
    all_docs_set = set(range(1, total_docs + 1)) # Create a set of all document  
    IDs
```

```
    return list(all_docs_set.difference(operand_set)) # Return documents not in
the operand set
```

```
# Example queries
```

```
query1 = ["apple", "banana"] # Query for documents containing both "apple"
and "banana"
```

```
query2 = ["apple", "orange"] # Query for documents containing "apple" or
"orange"
```

```
# Performing Boolean Model queries using inverted index
```

```
result1 = boolean_and(query1, inverted_index) # Get documents containing
both terms
```

```
result2 = boolean_or(query2, inverted_index) # Get documents containing
either of the terms
```

```
result3 = boolean_not("orange", inverted_index, len(documents)) # Get
documents not containing "orange"
```

```
# Printing results
```

```
print("Documents containing 'apple' and 'banana':", result1)
```

```
print("Documents containing 'apple' or 'orange':", result2)
```

```
print("Documents not containing 'orange':", result3)
```

```
print("Performed by 740_Pallavi & 743_Deepak")
```

Output:

```
Documents containing 'apple' and 'banana': [1, 2]
```

```
Documents containing 'apple' or 'orange': [1, 2, 3, 4]
```

```
Documents not containing 'orange': [2, 4]
```

```
Performed by 740_Pallavi & 743_Deepak
```

B) Implement the vector space model with TF-IDF weighting and cosine similarity:

Input:

```
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
# Import necessary libraries
import nltk # Import NLTK to download stopwords
from nltk.corpus import stopwords # Import stopwords from NLTK
import numpy as np # Import NumPy library
from numpy.linalg import norm # Import norm function from NumPy's linear algebra module
```

```
# Define the training and test sets of text documents
train_set = ["The sky is blue.", "The sun is bright."] # Documents
test_set = ["The sun in the sky is bright."] # Query
```

```
# Get the stopwords for English language from NLTK
nltk.download('stopwords')
stopWords = stopwords.words('english')
```

```
# Initialize CountVectorizer and TfidfTransformer objects
vectorizer = CountVectorizer(stop_words=stopWords) # CountVectorizer to
convert text to matrix of token counts
transformer = TfidfTransformer() # TfidfTransformer to convert matrix of
token counts to TF-IDF representation
```

```
# Convert the training and test sets to arrays of TF-IDF features
trainVectorizerArray = vectorizer.fit_transform(train_set).toarray() # Fit-
transform training set
testVectorizerArray = vectorizer.transform(test_set).toarray() # Transform test
set
```

```
# Display the TF-IDF arrays for training and test sets
```



```
print('Fit Vectorizer to train set', trainVectorizerArray)
print('Transform Vectorizer to test set', testVectorizerArray)

# Define a lambda function to calculate cosine similarity between vectors
cx = lambda a, b: round(np.inner(a, b) / (norm(a) * norm(b)), 3)

# Iterate through each vector in the training set
for vector in trainVectorizerArray:
    print(vector) # Display each vector in the training set
    # Iterate through each vector in the test set
    for testV in testVectorizerArray:
        print(testV) # Display each vector in the test set
        cosine = cx(vector, testV) # Calculate cosine similarity between vectors
        print(cosine) # Display the cosine similarity

# Fit the transformer to the training set and transform it to TF-IDF
representation
transformer.fit(trainVectorizerArray)
print()
print(transformer.transform(trainVectorizerArray).toarray())

# Fit the transformer to the test set and transform it to TF-IDF representation
transformer.fit(testVectorizerArray)
print()
tfidf = transformer.transform(testVectorizerArray)
print(tfidf.todense())
```

Output:

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\deepa\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
Fit Vectorizer to train set [[1 0 1 0]
[0 1 0 1]]
Transform Vectorizer to test set [[0 1 1 1]]
[1 0 1 0]
[0 1 1 1]
0.408
[0 1 0 1]
[0 1 1 1]
0.816

[[0.70710678 0.          0.70710678 0.          ]
[0.          0.70710678 0.          0.70710678]]

[[0.          0.57735027 0.57735027 0.57735027]]
Performed by 740_Pallavi & 743_Deepak
```

6/2/2024 TYCS Absent students
DUBEY SWATI RAJMANI
GOMANE SANKET SAKHARAM
NISHAD AMAN KUMAR
PATEL RAVIROUSHAN SURESH
SINGH ABHISHEK RAMESH
MALCHE ROSHANI DEVIDAS
ARULKAR TUSHAR ASHOK
PANDEY SHRIDHAR SACHIDANAND

Practical No: 3

Aim: Spelling Correction in IR Systems

- Develop a spelling correction module using edit distance algorithms.
- Integrate the spelling correction module into an information retrieval system.

Code:

Input:

A) Develop a spelling correction module using edit distance algorithms.

```
# A Naive recursive Python program to find the minimum number of operations  
# required to convert str1 to str2
```

```
def editDistance(str1, str2, m, n):  
    # If the first string is empty, insert all characters of the second string  
    if m == 0:  
        return n  
  
    # If the second string is empty, remove all characters of the first string  
    if n == 0:  
        return m  
  
    # If last characters are the same, ignore and recurse for remaining strings  
    if str1[m - 1] == str2[n - 1]:  
        return editDistance(str1, str2, m - 1, n - 1)  
  
    # If last characters are different, consider all three operations:  
    return 1 + min(  
        editDistance(str1, str2, m, n - 1), # Insert  
        editDistance(str1, str2, m - 1, n), # Remove  
        editDistance(str1, str2, m - 1, n - 1) # Replace  
    )
```

```
# Driver code  
str1 = "sunday"  
str2 = "saturday"
```

```
# Compute and print the edit distance  
print("Edit Distance is:", editDistance(str1, str2, len(str1), len(str2)))
```

Output:

```
----- RESTART -----  
>> Edit Distance is: 3  
|
```

Practical No: 4

Aim: Evaluation Metrics for IR Systems

- A) Calculate precision, recall, and F-measure for a given set of retrieval results.
- B) Use an evaluation toolkit to measure average precision and other evaluation metrics.

Code:

Input:

- A) Calculate precision, recall, and F-measure for a given set of retrieval results.

```
def calculate_metrics(retrieved_set, relevant_set):
    true_positive = len(retrieved_set.intersection(relevant_set))
    false_positive = len(retrieved_set.difference(relevant_set))
    false_negative = len(relevant_set.difference(retrieved_set))

    """
    (Optional) PPT values:
    true_positive = 20
    false_positive = 10
    false_negative = 30
    """

    print("True Positive:", true_positive,
          "\nFalse Positive:", false_positive,
          "\nFalse Negative:", false_negative, "\n")

    precision = true_positive / (true_positive + false_positive) if (true_positive +
false_positive) != 0 else 0
    recall = true_positive / (true_positive + false_negative) if (true_positive +
false_negative) != 0 else 0
    f_measure = (2 * precision * recall / (precision + recall)) if (precision + recall) != 0
else 0

    return precision, recall, f_measure
```

```
retrieved_set = set(["doc1", "doc2", "doc3"]) # Predicted set
relevant_set = set(["doc1", "doc4"]) # Actually Needed set (Relevant)

precision, recall, f_measure = calculate_metrics(retrieved_set, relevant_set)

print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F-measure: {f_measure}")
```

Output:

```
True Positive: 1
False Positive: 2
False Negative: 1

Precision: 0.3333333333333333
Recall: 0.5
F-measure: 0.4
>>> |
```

B) Use an evaluation toolkit to measure average precision and other evaluation metrics.

Input:

```
from sklearn.metrics import average_precision_score

y_true = [0, 1, 1, 0, 1, 1] # Binary Prediction
y_scores = [0.1, 0.4, 0.35, 0.8, 0.65, 0.9] # Model's estimation score

average_precision = average_precision_score(y_true, y_scores)

print(f'Average precision-recall score: {average_precision}')
```

Output:

Result

0.8041666666666667

Practical No: 5

Aim: Text Categorization

- A) Implement a text classification algorithm (e.g., Naive Bayes or Support Vector Machines).**
- B) Train the classifier on a labelled dataset and evaluate its performance.**

Code:

Input:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report

# Load the CSV file
df = pd.read_csv(r"C:\Users\Administrator\Documents\Sem 6\IR\Dataset.csv")

# Combining the relevant columns into a single feature
data = df["covid"] + " " + df["fever"]
X = data.astype(str) # Convert to string for text processing
y = df['flu'] # Labels

# Splitting the data into training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Converting text data into a bag-of-words format
vectorizer = CountVectorizer()

# Fit and transform the training data
X_train_counts = vectorizer.fit_transform(X_train)

# Transform the test data
X_test_counts = vectorizer.transform(X_test)
```



```
# Initializing and training the Naive Bayes classifier
classifier = MultinomialNB()
classifier.fit(X_train_counts, y_train)

# Loading another dataset to test if the model is working properly
data1 = pd.read_csv(r"C:\Users\Administrator\Documents\Sem 6\IR\Test.csv")
new_data = data1["covid"] + " " + data1["fever"]

# Convert new data using the trained vectorizer
new_data_counts = vectorizer.transform(new_data.astype(str))

# Predicting results for the new dataset
predictions = classifier.predict(new_data_counts)

# Output the results
print("Predictions for new dataset:")
print(predictions)

# Retrieving the accuracy and classification report
accuracy = accuracy_score(y_test, classifier.predict(X_test_counts))
print(f"\nAccuracy: {accuracy:.2f}")

print("\nClassification Report:")
print(classification_report(y_test, classifier.predict(X_test_counts)))

# Convert the predictions to a DataFrame
predictions_df = pd.DataFrame(predictions, columns=['flu_prediction'])

# Concatenate the original DataFrame with the predictions DataFrame
data1 = pd.concat([data1, predictions_df], axis=1)

# Write the DataFrame back to CSV
data1.to_csv(r"C:\Users\Administrator\Documents\Sem 6\IR\Test1.csv",
index=False)
```

Output:

- Documents/Sem 6/IR/prac_5.py"

```
['yes' 'no' 'yes' 'no' 'no' 'yes' 'no' 'yes' 'no' 'yes']
```

Accuracy: 1.00

Classification Report:

	precision	recall	f1-score	support
no	1.00	1.00	1.00	2
accuracy			1.00	2
macro avg	1.00	1.00	1.00	2
weighted avg	1.00	1.00	1.00	2

Performed by 740_Pallavi & 743_Deepak

[illegible]

Practical No: 6

Aim: Clustering for Information Retrieval

- Implement a clustering algorithm (e.g., K-means or hierarchical clustering).
- Apply the clustering algorithm to a set of documents and evaluate the clustering results.

Code:

Input:

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.cluster import KMeans
```

```
documents = ["Cats are known for their agility and grace", #cat doc1
             "Dogs are often called 'man's best friend'.", #dog doc1
             "Some dogs are trained to assist people with disabilities.", #dog doc2
             "The sun rises in the east and sets in the west.", #sun doc1
             "Many cats enjoy climbing trees and chasing toys.", #cat doc2
             ]
# Create a TfidfVectorizer object
vectorizer = TfidfVectorizer(stop_words='english')

# Learn vocabulary and idf from training set.
X = vectorizer.fit_transform(documents)

# Perform k-means clustering
kmeans = KMeans(n_clusters=3, random_state=0).fit(X)

# Print cluster labels for each document
print(kmeans.labels_)
```

Output:

```
[0 1 1 2 0]
Performed by 740_Pallavi & 743_Deepak
```

Practical No: 7

Aim: Web Crawling and Indexing

A) Develop a web crawler to fetch and index web pages.

B) Handle challenges such as robots.txt, dynamic content, and crawling delays.

Code:

Input:

```
import requests
from bs4 import BeautifulSoup
import time
from urllib.parse import urljoin, urlparse
from urllib.robotparser import RobotFileParser

def get_html(url):
    headers = {'User-Agent': 'Mozilla/5.0 (Windows NT 10.0; Win64; x64)
AppleWebKit/537.36 (KHTML, like Gecko) Chrome/58.0.3029.110
Safari/537.3'}
    try:
        response = requests.get(url, headers=headers)
        response.raise_for_status()
        return response.text
    except requests.exceptions.HTTPError as errh:
        print(f"HTTP Error: {errh}")
    except requests.exceptions.RequestException as err:
        print(f"Request Error: {err}")
    return None

def save_robots_txt(url):
    try:
        robots_url = urljoin(url, '/robots.txt')
        robots_content = get_html(robots_url)
        if robots_content:
            with open('robots.txt', 'wb') as file:
```

```
        file.write(robots_content.encode('utf-8-sig'))
except Exception as e:
    print(f"Error saving robots.txt: {e}")

def load_robots_txt():
    try:
        with open('robots.txt', 'rb') as file:
            return file.read().decode('utf-8-sig')
    except FileNotFoundError:
        return None

def extract_links(html, base_url):
    soup = BeautifulSoup(html, 'html.parser')
    links = []
    for link in soup.find_all('a', href=True):
        absolute_url = urljoin(base_url, link['href'])
        links.append(absolute_url)
    return links

def is_allowed_by_robots(url, robots_content):
    parser = RobotFileParser()
    parser.parse(robots_content.split('\n'))
    return parser.can_fetch('*', url)

def crawl(start_url, max_depth=3, delay=1):
```

```
visited_urls = set()

def recursive_crawl(url, depth, robots_content):
    if depth > max_depth or url in visited_urls or not
is_allowed_by_robots(url, robots_content):
        return
    visited_urls.add(url)
    time.sleep(delay)

    html = get_html(url)
    if html:
        print(f"Crawling {url}")
        links = extract_links(html, url)

        for link in links:
            recursive_crawl(link, depth + 1, robots_content)

save_robots_txt(start_url)
robots_content = load_robots_txt()
if not robots_content:
    print("Unable to retrieve robots.txt. Crawling without restrictions.")

recursive_crawl(start_url, 1, robots_content)

# Example usage:
print("Performed by 740_Pallavi & 743_Deepak")
crawl('https://wikipedia.com', max_depth=2, delay=2)
```

Output:

Performed by 740_Pallavi & 743_Deepak
Crawling <https://wikipedia.com>
Crawling <https://en.wikipedia.org/>
Crawling <https://ja.wikipedia.org/>
Crawling <https://ru.wikipedia.org/>
Crawling <https://de.wikipedia.org/>
Crawling <https://es.wikipedia.org/>
Crawling <https://fr.wikipedia.org/>
Crawling <https://it.wikipedia.org/>
Crawling <https://zh.wikipedia.org/>
Crawling <https://fa.wikipedia.org/>
Crawling <https://pl.wikipedia.org/>
Crawling <https://ar.wikipedia.org/>

robot.txt file:

```
robots.txt x
robots.txt
1  # robots.txt for http://www.wikipedia.org/ and friends
2  #
3  # Please note: There are a lot of pages on this site, and there are
4  # some misbehaved spiders out there that go _way_ too fast. If you're
5  # irresponsible, your access to the site may be blocked.
6  #
7
8  # Observed spamming large amounts of https://en.wikipedia.org/?curid=NNNNNN
9  # and ignoring 429 ratelimit responses, claims to respect robots:
10 # http://mj12bot.com/
11 User-agent: MJ12bot
12 Disallow: /
13
14 # advertising-related bots:
15 User-agent: Mediapartners-Google*
16 Disallow: /
17
18 # Wikipedia work bots:
19 User-agent: IsraBot
20 Disallow:
21
22 User-agent: Orthogaffe
23 Disallow:
24
25 # Crawlers that are kind enough to obey, but which we'd rather not have
26 # unless they're feeding search engines.
27 User-agent: UbiCrawler
28 Disallow: /
29
30 User-agent: DOC
31 Disallow: /
```

Practical No: 8

Aim: Link Analysis and PageRank

- A) Implement the PageRank algorithm to rank web pages based on link analysis.
- B) Apply the PageRank algorithm to a small web graph and analyze the results.

Code:

Input:

```
import numpy as np

def page_rank(graph, damping_factor=0.85, max_iterations=100,
              tolerance=1e-6):
    # Get the number of nodes
    num_nodes = len(graph)

    # Initialize PageRank values
    page_ranks = np.ones(num_nodes) / num_nodes

    # Iterative PageRank calculation
    for _ in range(max_iterations):
        prev_page_ranks = np.copy(page_ranks)

        for node in range(num_nodes):
            # Calculate the contribution from incoming links
            incoming_links = [i for i, v in enumerate(graph) if node in v]

            if not incoming_links:
                page_ranks[node] = (1 - damping_factor) / num_nodes
            else:
                page_ranks[node] = (1 - damping_factor) / num_nodes + \
                    damping_factor * sum(prev_page_ranks[link] /
                    len(graph[link]) for link in incoming_links)

        # Check for convergence
        if np.linalg.norm(page_ranks - prev_page_ranks, 2) < tolerance:
            break

    return page_ranks
```



```
# Example usage
if __name__ == "__main__":
    # Define a simple directed graph as an adjacency list
    # Each index represents a node, and the list at that index contains nodes
    # to which it has outgoing links
    web_graph = [
        [1, 2], # Node 0 has links to Node 1 and Node 2
        [0, 2], # Node 1 has links to Node 0 and Node 2
        [0, 1], # Node 2 has links to Node 0 and Node 1
        [1, 2], # Node 3 has links to Node 1 and Node 2
    ]

    # Calculate PageRank
    result = page_rank(web_graph)

    # Display PageRank values
    for i, pr in enumerate(result):
        print(f"Page {i}: {pr:.4f}")
```

Output:

```
Page 0: 0.3134
Page 1: 0.3246
Page 2: 0.3246
Page 3: 0.0375
>>> |
```